


Lecture 3

Processes and Filters

Kernel Data Structures

- Information about each process.
- **Process table:** contains an entry for every process in the system.
- **Open-file table:** contains at least one entry for every open file in the system.



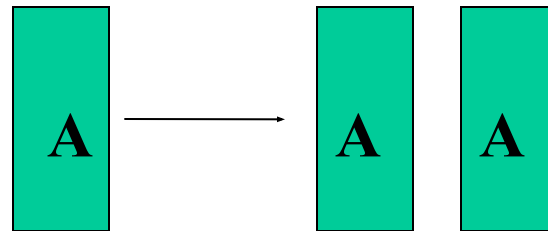
Unix Processes

Process: An entity of execution

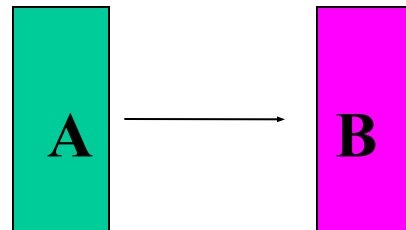
- *Definitions*
 - **program**: collection of bytes stored in a file that can be run
 - **image**: computer execution environment of program
 - **process**: execution of an image
- Unix can execute many processes simultaneously.

Process Creation

- Interesting trait of UNIX
- **fork** system call clones the current process



- **exec** system call reinitializes current process with new image



- A **fork** is typically followed by an **exec**

Process Setup

- All of the per process information is copied with the **fork** operation
 - Working directory
 - Open files
- *Copy-on-write* makes this efficient
- Before **exec**, these values can be modified
- Both processes receive return from **fork()**
 - Child gets 0; parent gets PID of child

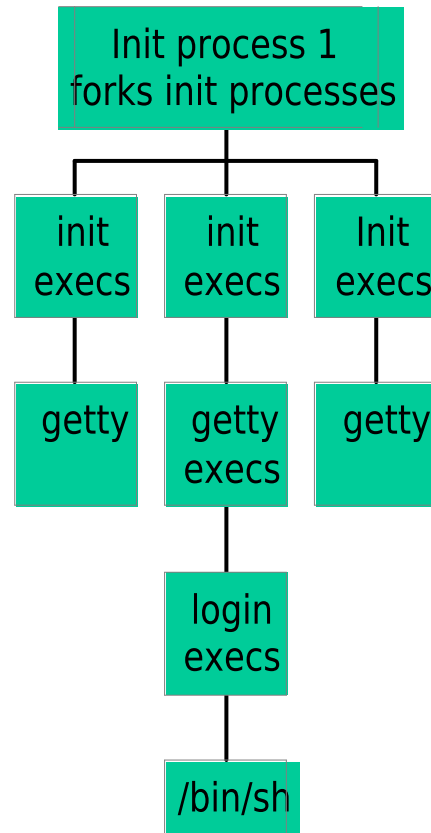
fork and exec

Example: the shell

```
while(1) {
    display_prompt();
    read_input(cmd, params);
    pid = fork();          /* create child */
    if (pid != 0)
        waitpid(-1, &stat, 0); /* parent waits */
    else
        execve(cmd, params, 0); /* child execs */
}
```

Unix process genealogy

Process generation



Background Jobs

- By default, executing a command in the shell will wait for it to exit before printing out the next prompt
- Trailing a command with `&` allows the shell and command to run simultaneously

```
$ /bin/sleep 10 &  
[1] 3424  
$
```


Program Arguments

- When a process is started, it is sent a list of strings
 - **argv**, **argc**
- The process can use this list however it wants to

```
find . -name 'foo*'
```

0	find
1	.
2	-name
3	foo*

Ending a process

- When a process ends, there is a return code associated with the process
- This is a positive integer:
 - 0 means success
 - > 0 represents various kinds of failure.
Up to process to decide meaning.

Process Information Maintained

- Working directory
- File descriptor table
- Process id
 - number used to identify process
- Process group id
 - number used to identify set of processes
- Parent process id
 - process id of the process that created the process

Process Information Maintained

- Umask
 - Default file permissions for new file

We haven't talked about these yet:

- Effective user and group id
 - The user and group this process is running with permissions as
- Real user and group id
 - The user and group that invoked the process
- Environment variables

Setuid and Setgid Mechanisms

- The kernel can set the effective user and group ids of a process to something different than the real user and group
 - Files executed with a setuid or setgid flag set cause these values to change
- Make it possible to do privileged tasks:
 - Change your password
- Open up a can of worms for security if buggy

Environment of a Process

- A set of name-value pairs associated with a process
- Keys and values are strings
- Passed to children processes
- Cannot be passed back up
- Common examples:
 - **PATH**: Where to search for programs
 - **TERM**: Terminal type



The PATH environment variable

- Colon-separated list of directories.
- Non-absolute pathnames of executables are only executed if found in the list.
 - Searched left to right

- Example:

```
$ myprogram
```

```
sh: myprogram not found
```

```
$ PATH=/bin:/usr/bin:/home/kornj/bin
```

```
$ myprogram
```

```
hello!
```



Shell Variables

- Shells have several mechanisms for creating variables. A variable is a name representing a string value. Example: **PATH**
 - Shell variables can save time and reduce typing errors
- Allows you to store and manipulate information
 - Eg: `ls $DIR > $FILE`
- Two types: **local** and **environment**
 - *local* are set by the user or by the shell itself
 - *environment* come from the operating system and are passed to children

Variables (con't)

- Syntax varies by shell

`varname=value` # sh, ksh, bash

`set varname = value` # csh

- To access the value: `$varname`

- Turn local variable into environment:

`export varname` # sh, ksh, bash

`setenv varname value` # csh

Environmental Variables

NAME	MEANING
\$HOME	Absolute pathname of your home directory
\$PATH	A list of directories to search for
\$MAIL	Absolute pathname to mailbox
\$USER	Your user id
\$SHELL	Absolute pathname of login shell
\$TERM	Type of your terminal
\$PS1	Prompt


Inter-process Communication

Ways in which processes communicate:

- Passing arguments, environment
- Read/write regular files
- Exit values
- **Signals**
- **Pipes**

Signals

- **Signal:** A message a process can send to a process or process group, if it has appropriate permissions.
- Message type represented by a symbolic name
- For each signal, the **receiving process** can:
 - Explicitly ignore signal
 - Specify action to be taken upon receipt (**signal handler**)
 - Otherwise, default action takes place (usually process is killed)
- Common signals:
 - SIGKILL, SIGTERM, SIGINT
 - SIGSTOP, SIGCONT
 - SIGSEGV, SIGBUS



An Example of Signals

- When a child exits, it sends a **SIGCHLD** signal to its parent.
- If a parent wants to be notified that a child has exited, it tells the system it wants to catch the **SIGCHLD** signal, via **waitpid()**
- Default action for **SIGCHLD** signal is to ignore it



Process Subsystem utilities

ps	monitors status of processes
kill	send a signal to a pid
wait	parent process wait for one of its children to terminate
nohup	makes a command immune to the hangup and terminate signal
sleep	sleep in seconds
nice	run processes at low priority

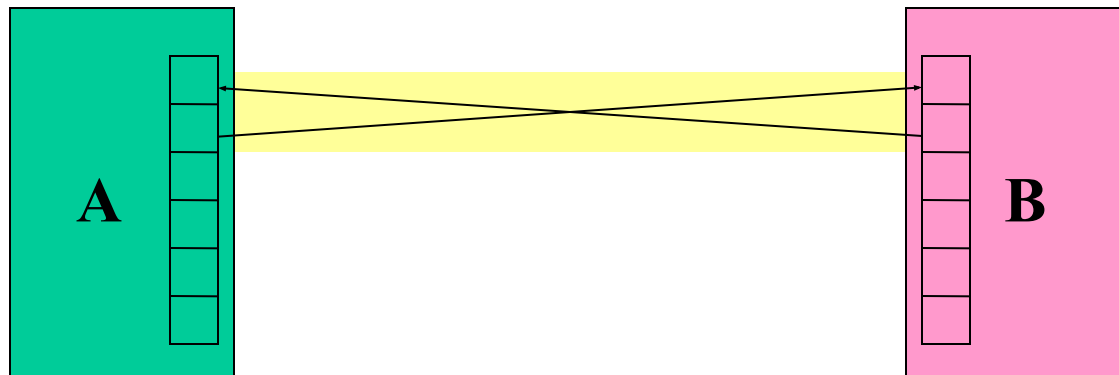
Pipes



One of the cornerstones of UNIX

Pipes

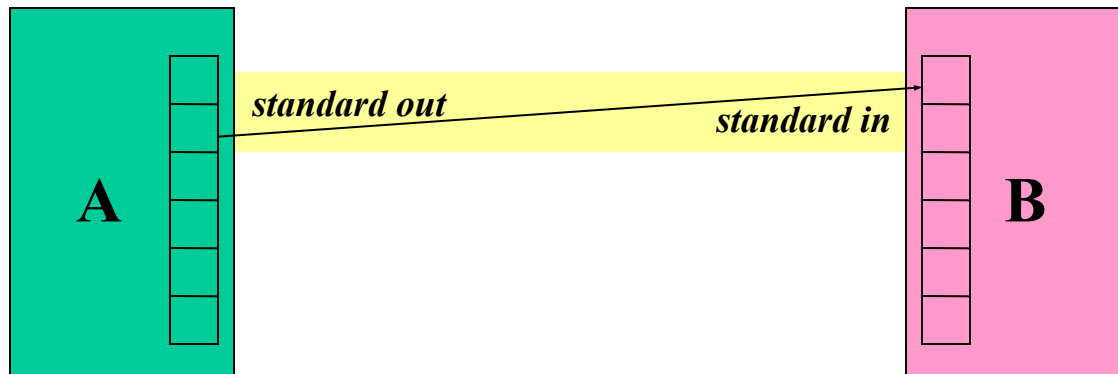
- General idea: The input of one program is the output of the other, and vice versa



- Both programs run at the same time

Pipes (2)

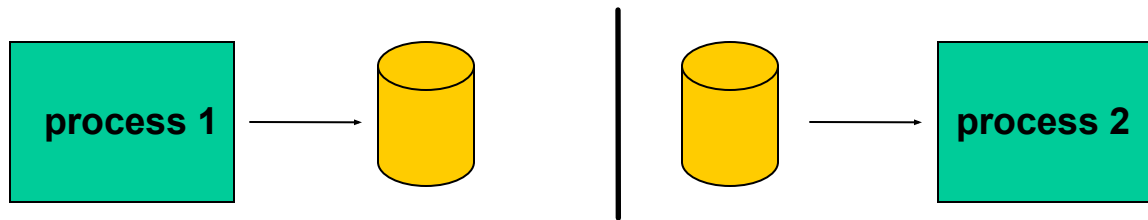
- Often, only one end of the pipe is used



- Could this be done with files?

File Approach

- Run first program, save output into file
- Run second program, using file as input



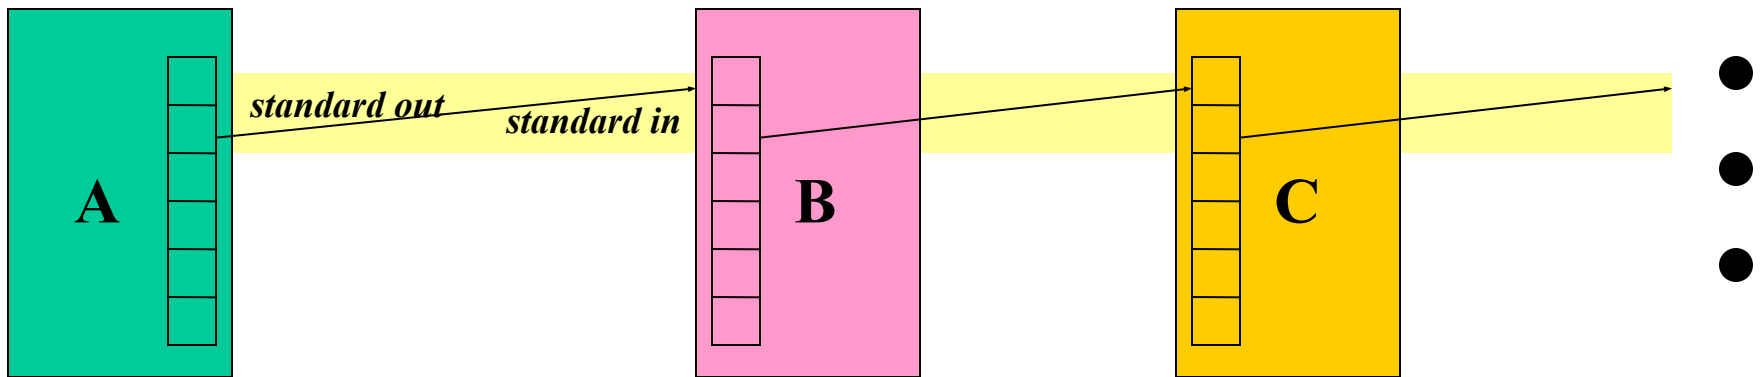
- Unnecessary use of the disk
 - Slower
 - Can take up a lot of space
- Makes no use of multi-tasking

More about pipes

- When process tries to read from pipe but nothing is available...
 - UNIX puts the reader to sleep until data available
- When reader cannot keep up with writer...
 - Unread data is buffered up to the *pipe size*
 - If pipe fills up, writer put to sleep until the reader frees up space (by reading)
- Multiple readers and writers possible with pipes.


More about Pipes

- Pipes are often chained together
 - Called *filters*



Interprocess Communication For Unrelated Processes

- FIFO (*named pipes*)
 - A special file that when opened represents pipe
- System V IPC
 - message queues
 - semaphores
 - shared memory
- Sockets (client/server model)




Ceci, n'est pas une pipe.

Pipelines

- Output of one program becomes input to another
 - Uses concept of UNIX **pipes**
- Example: `$ who | wc -l`
 - counts the number of users logged in
- Pipelines can be long

`filter1 | filter2 | filter3 | ... | filterN`



What's the difference?

Both of these commands send input to ***command*** from a file instead of the terminal:

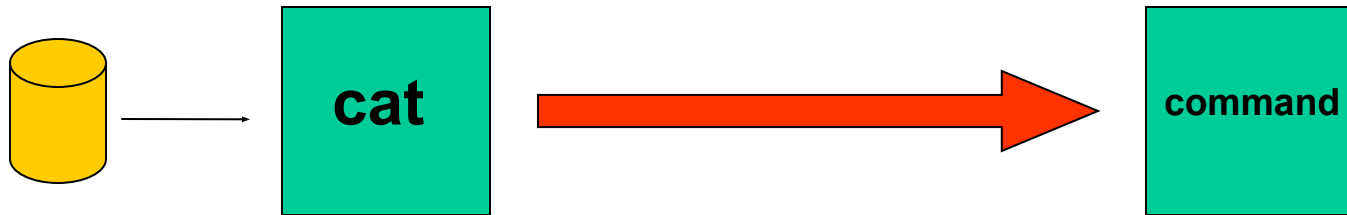
```
$ cat file | command
```

vs.

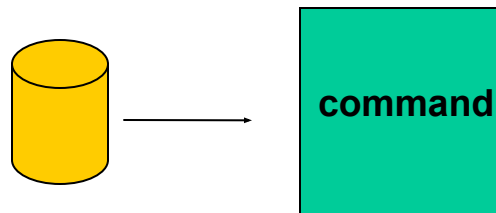
```
$ command < file
```

An Extra Process

```
$ cat file | command
```

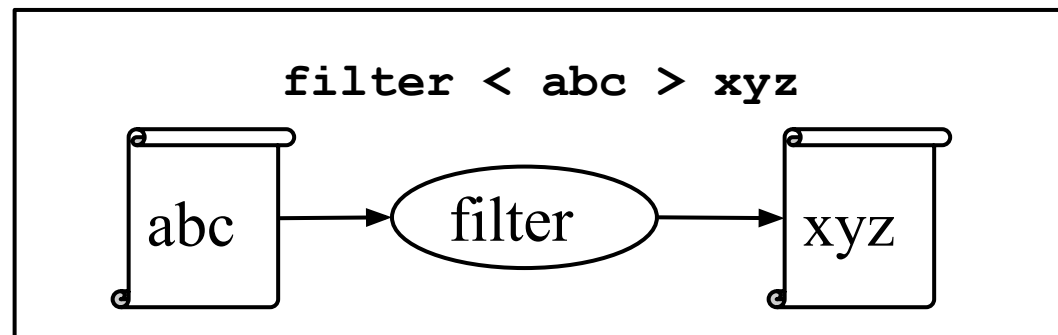


```
$ command < file
```



Introduction to Filters

- A class of Unix tools called *filters*.
 - Utilities that read from standard input, transform the file, and write to standard out
- Using filters can be thought of as *data oriented programming*.
 - Each step of the computation transforms data *stream*.



Examples of Filters

- **sort**
 - Input: lines from file/stdin
 - Output: lines sorted
- **grep**
 - Input: lines from file/stdin
 - Output: lines that match pattern as argument
- **awk**
 - Programmable filter

cat: The simplest filter

- The **cat** command copies its input to output unchanged (*identity filter*). When supplied a list of file names, it concatenates them onto stdout.
- Some options:
 - n number output lines (starting from 1)
 - v display control-characters in visible form (e.g. ^C)

```
cat file*
```

```
ls | cat -n
```

head

- Display the first few lines of a specified file
- Syntax: *head [-n] [filename...]*
 - n* - number of lines to display, default is 10
 - filename...* - list of filenames to display
- When more than one filename is specified, the start of each files listing displays
 - `==> filename <==`

tail

- Displays the last part of a file
- Syntax: *tail +|-number [lbc] [f] [filename]*
or: *tail +|-number [l] [rf] [filename]*
 - +number* - begins copying at distance *number* from beginning of file, if *number* isn't given, defaults to 10
 - number* - begins from end of file
 - l,b,c* - *number* is in units of lines/block/characters
 - r* - print in reverse order (lines only)
 - f* - if input is not a pipe, do not terminate after end of file has been copied but loop. This is useful to monitor a file being written by another process

head and tail examples

```
head /etc/passwd
```

```
head *.c
```

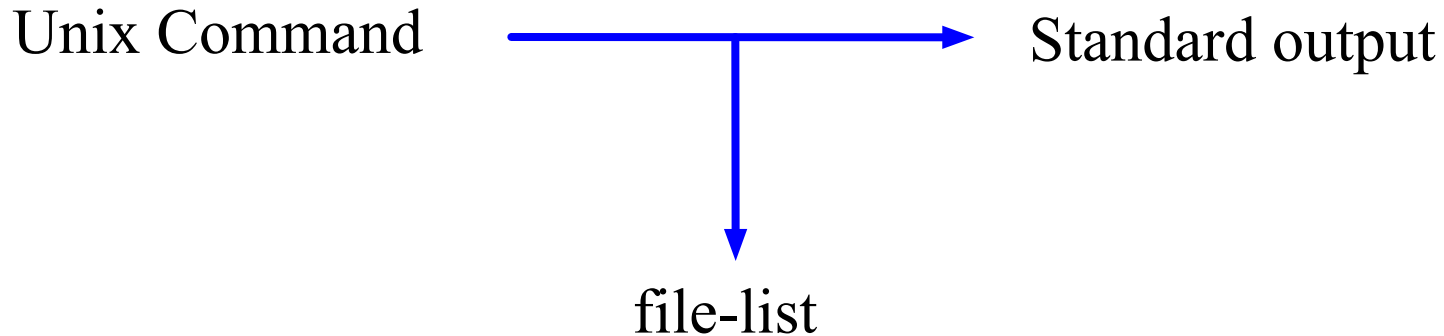
```
tail +20 /etc/passwd
```

```
ls -lt | tail -3
```


```
head -100 /etc/passwd | tail -5
```

```
tail -f /var/httpd/access_log
```

tee



- Copy standard input to standard output and one or more files
 - Captures intermediate results from a filter in the pipeline



tee (2)

- Syntax: *tee* [*-ai*] *file-list*
 - a* - append to output file rather than overwrite, default is to overwrite (replace) the output file
 - i* - ignore interrupts
 - file-list* - one or more file names for capturing output
- Examples:

```
head -10 | tee first_10 | tail -5  
who | tee user_list | wc -l
```


Unix Text Files: Delimited Data

Tab Separated

John	99
Anne	75
Andrew	50
Tim	95
Arun	33
Sowmya	76

Pipe-separated

COMP1011 2252424 Abbot, Andrew John 3727 1 M
COMP2011 2211222 Abdurjh, Saeed 3640 2 M
COMP1011 2250631 Accent, Aac-Ek-Murhg 3640 1 M
COMP1021 2250127 Addison, Blair 3971 1 F
COMP4012 2190705 Allen, David Peter 3645 4 M
COMP4910 2190705 Allen, David Pater 3645 4 M

Colon-separated

root:ZHolHAHZw8As2:0:0:root:/root:/bin/ksh
jas:nJz3ru5a/44Ko:100:100:John Shepherd:/home/jas:/bin/ksh
cs1021:iZ3sO90O5eZY6:101:101:COMP1021:/home/cs1021:/bin/bash
cs2041:rX9KwSSPqkLyA:102:102:COMP2041:/home/cs2041:/bin/csh
cs3311:mLRiCIvmtI9O2:103:103:COMP3311:/home/cs3311:/bin/sh

cut: select columns

- The **cut** command prints selected parts of input lines.
 - can select columns (assumes tab-separated input)
 - can select a range of character positions
- Some options:
 - f *listOfCols*: print only the specified columns (tab-separated) on output
 - c *listOfPos*: print only chars in the specified positions
 - d *c*: use character *c* as the column separator
- Lists are specified as ranges (e.g. 1–5) or comma-separated (e.g. 2, 4, 5).

cut examples

```
cut -f 1 < data
```

```
cut -f 1-3 < data
```

```
cut -f 1,4 < data
```

```
cut -f 4- < data
```

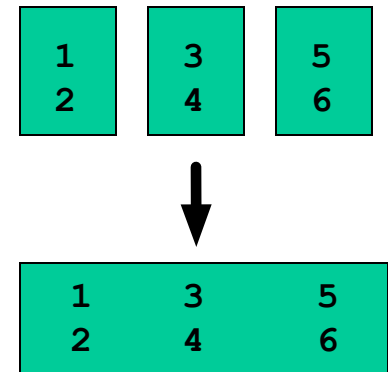
```
cut -d'|' -f 1-3 < data
```

```
cut -c 1-4 < data
```

Unfortunately, there's no way to refer to "last column" without counting the columns.

paste: join columns

- The **paste** command displays several text files "in parallel" on output.
- If the inputs are files **a**, **b**, **c**
 - the first line of output is composed of the first lines of **a**, **b**, **c**
 - the second line of output is composed of the second lines of **a**, **b**, **c**
- Lines from each file are separated by a tab character.
- If files are different lengths, output has all lines from longest file, with empty strings for missing



paste example

```
cut -f 1 < data > data1
```

```
cut -f 2 < data > data2
```

```
cut -f 3 < data > data3
```

```
paste data1 data3 data2 > newdata
```

sort: Sort lines of a file

- The **sort** command copies input to output but ensures that the output is arranged in ascending order of lines.
 - By default, sorting is based on ASCII comparisons of the whole line.
- Other features of **sort**:
 - understands text data that occurs in columns.
(can also sort on a column other than the first)
 - can distinguish numbers and sort appropriately
 - can sort files *in place* as well as behaving like a filter

sort: Options

- Syntax: *sort [-dftnr] [-o filename] [filename(s)]*
 - d* Dictionary order, only letters, digits, and whitespace are significant in determining sort order
 - f* Ignore case (fold into lower case)
 - t* Specify delimiter
 - n* Numeric order, sort by arithmetic value instead of first digit
 - r* Sort in reverse order
 - o filename* - write output to filename, filename can be the same as one of the input files
- Lots of more options...

sort: Specifying fields

- Delimiter : **-t***d*
- Ancient way:
 - **+f[.c][options] [-f[.c][options]**
+1.1 -2 +0 -2 +3n
 - Exclusive
 - Start from 0 (unlike cut, which starts at 1)
- Modern way:
 - **-k f[.c][options] [,f[.c][options]]**
-k2.2,2 -k1,2 -k4n
 - Inclusive
 - Start from 1

sort Examples

```
sort +2nr < data
```

```
sort -k2nr data
```

```
sort -t: -k5 /etc/passwd
```

```
sort -o mydata mydata
```

uniq: list **unique** items

- Remove or report adjacent duplicate lines
- Syntax: *uniq* [*-cdu*] [*input-file*] [*output-file*]
 - **c** Supersede the *-u* and *-d* options and generate an output report with each line preceded by an occurrence count
 - **d** Write only the duplicated lines
 - **u** Write only those lines which are not duplicated
- The default output is the union (combination) of *-d* and *-u*

wc: Counting results

- The word count utility, **wc**, counts the number of lines, characters or words
- Options:
 - l** Count lines
 - w** Count words
 - c** Count characters
- Default: count lines, words and chars

wc and uniq Examples

```
who | sort | uniq -d
```

```
wc my_essay
```

```
who | wc
```

```
sort file | uniq | wc -l
```

```
sort file | uniq -d | wc -l
```

```
sort file | uniq -u | wc -l
```

tr: **tr**anslate characters

- Copies standard input to standard output with substitution or deletion of selected characters
- Syntax: *tr* [*-cds*] [*string1*] [*string2*]
 - d** delete all input characters contained in *string1*
 - c** complements the characters in *string1* with respect to the entire ASCII character set
 - s** squeeze all strings of repeated output characters in the last operand to single characters

tr (continued)

- *tr* reads from standard input (only).
 - Any character that does not match a character in *string1* is passed to *standard output* unchanged
 - Any character that does match a character in *string1* is translated into the corresponding character in *string2* and then passed to *standard output*

- Examples

<code>tr s z</code>	replaces all instances of <i>s</i> with <i>z</i>
<code>tr so zx</code>	replaces all instances of <i>s</i> with <i>z</i> and <i>o</i> with <i>x</i>
<code>tr a-z A-Z</code>	replaces all lower case characters with upper case characters
<code>tr -d a-c</code>	deletes all a-c characters

tr uses

- Change delimiter

```
tr '|' ':'
```

- Rewrite numbers

```
tr ,. .,
```

- Import DOS files

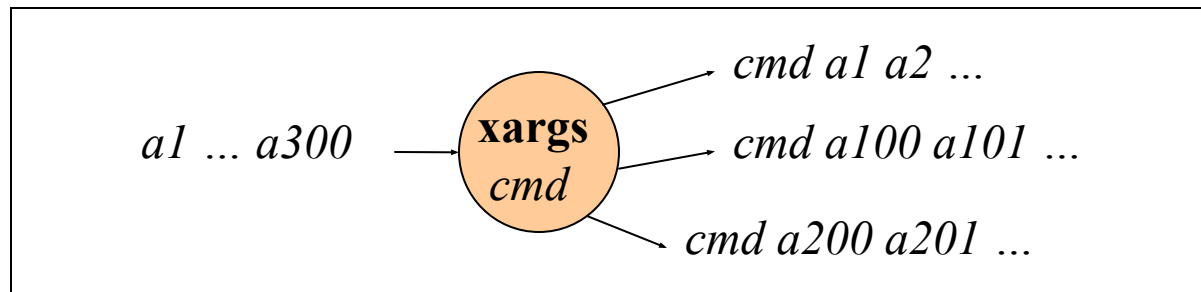
```
tr -d '\r' < dos_file
```

- Find printable ASCII in a binary file

```
tr -cd '\na-zA-Z0-9 ' < binary_file
```

xargs

- Unix limits the size of arguments and environment that can be passed down to child
- What happens when we have a list of 10,000 files to send to a command?
- **xargs** solves this problem
 - Reads arguments as standard input
 - Sends them to commands that take file lists
 - May invoke program several times depending on size of arguments



find utility and xargs

- `find . -type f -print | xargs wc -l`
 - `-type f` for files
 - `-print` to print them out
 - `xargs` invokes `wc` 1 or more times

```
wc -l a b c d e f g
wc -l h i j k l m n o
```

...

- Compare to:
`find . -type f -exec wc -l {} \;`

Next Time

- Regular Expressions
 - Allow you to search for text in files
 - **grep** command
- We will soon learn how to write *scripts* that use these utilities in interesting ways.